

又拍云 OpenResty/Nginx 服务优化实践

OpenResty Con 2018

张超 (Alex Zhang)





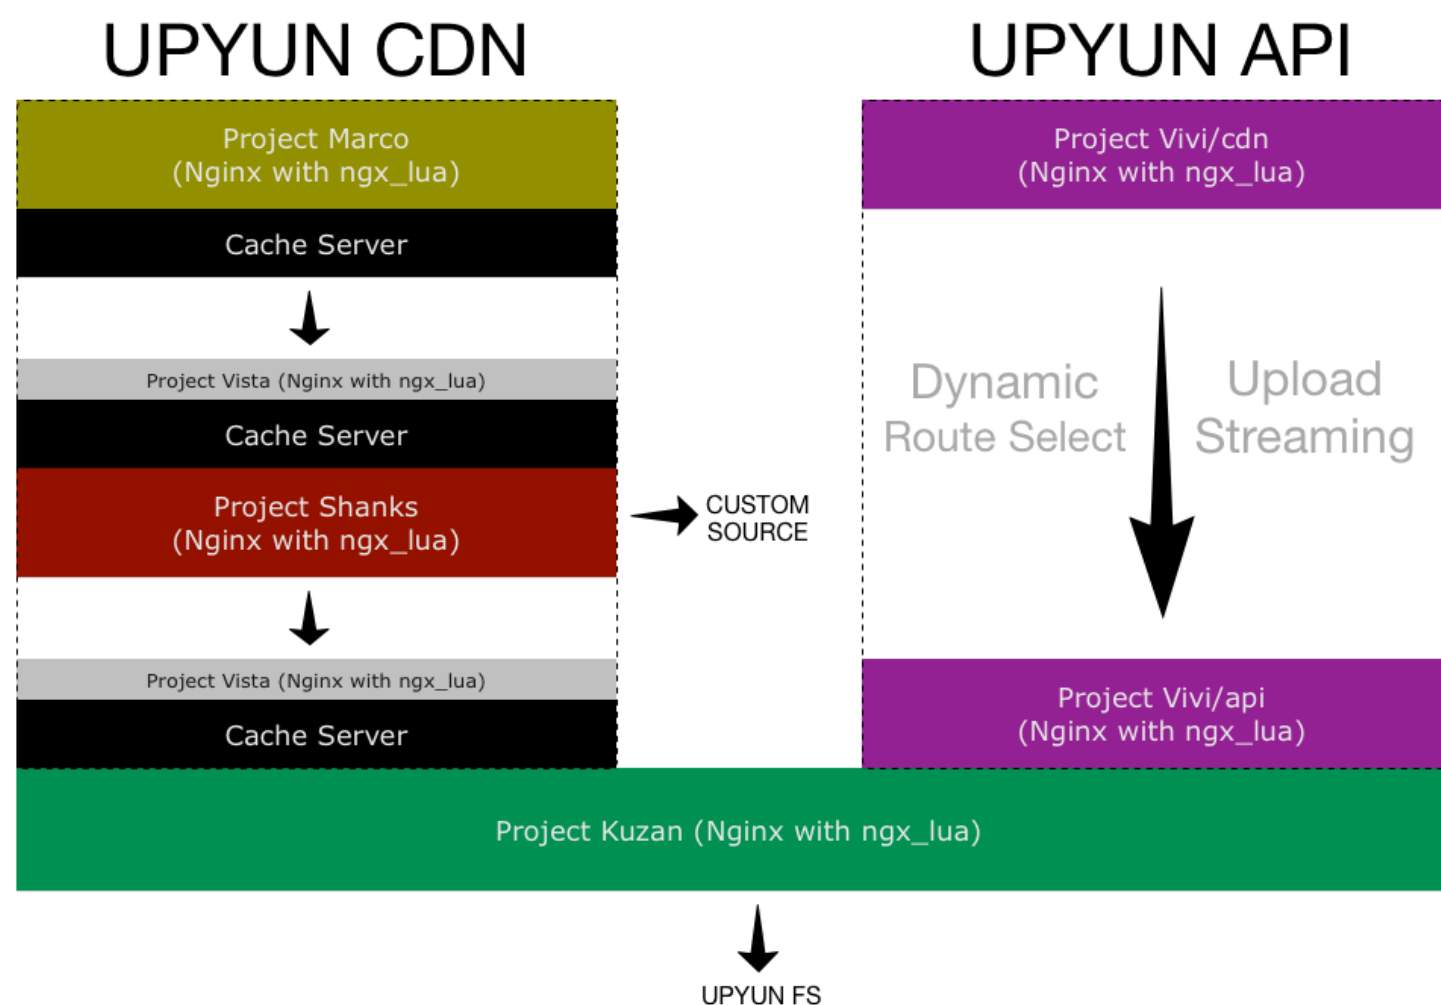
System Development Engineer at UPYUN Inc.

☆ Email: zchao1995@gmail.com

☆ Github: <https://github.com/tokers>

OpenResty in UPYUN

- ◉ Since 2013.
- ◉ All CDN services (except the Cache Server) are based on **OpenResty**.
- ◉ The gateway of UPYUN Storage System is based on **OpenResty**.
- ◉ All API services use **Kong** as the gateway.





<https://github.com/upyun/upyun-resty>

- Many lua-resty libraries
- Many Nginx C modules
- Reports and fixes some bugs about OpenResty and Nginx.

So we have entered a stage where we need to do the **deep optimization** for our services.



Performance Analysis & Tips

Name Resolution Management

SSL Acceleration in Practice

Resources analysis

- top
- pidstat
- vmstat
- iostat
- sar
-

Resource analysis focuses on resources' utilization (U), saturation (S) and errors (E).

Workload analysis

- perf
- systemtap
- bcc/ebpf
- flamegraph
-

Workload analysis focuses on application's own performance, like latency, bad status (5xx) ratio.

But sometimes tool chain is incomplete/absent
in the production environment.

So build a docker image with:

- compilers/linkers
- perf
- systemtap
- flamegraph scripts
- gdb
- mozilla rr
- valgrind
- ...

💡 Performance diagnosis in a privileged docker container.

```
docker run --it --rm --privileged \  
-v /lib/modules:/lib/modules:ro \  
-v /etc/localtime:/etc/localtime:ro \  
-v /usr/src:/usr/src:ro \  
--net=host \  
--pid=host upyun_stap:latest
```

ngx.ctx & ngx.var

Both `ngx.ctx` and `ngx.var` can be used to store request-specific data.

`ngx.ctx` is a Lua table while `ngx.var` is an interface to get/set Nginx's variables.

ngx.ctx is better

- Lua table is fast
- Return value of `ngx.var.*` can only be a string
- hash calculations, memory allocations inside `ngx.var.*` calls

💡 Cache `ngx.ctx` to avoid the expensive meta-method calls.

```
local function f()  
    local ctx = ngx.ctx  
    ctx.foo = "bar"  
    ctx.bar = "foo"  
    ctx.done = true  
end
```

The life cycle of `ngx.ctx` is limited to a Nginx location.

Data will be lost when Nginx internal redirect happens.

💡 Use `lua-resty-ctxdump` to bypass this limitation.

```
location /t1 {
    set $ctx_ref "";
    content_by_lua_block {
        local ctxdump = require "resty.ctxdump"
        ngx.ctx = {
            Date = "Wed May  3 15:18:04 CST 2017",
        }
        ngx.var.ctx_ref = ctxdump.stash_ngx_ctx()
        ngx.exec("/t2")
    }
}
location /t2 {
    internal;
    content_by_lua_block {
        local ctxdump = require "resty.ctxdump"
        ngx.ctx = ctxdump.apply_ngx_ctx(ngx.var.ctx_ref)
    }
}
```

HTTP headers

```
local user_agent = ngx.req.get_headers()["User-Agent"]  
local cookie = ngx.req.get_headers()["Cookie"]  
local x_forwarded_for = ngx.req.get_headers()["X-Forwarded-For"]
```

`ngx.req.get_headers()` will return **all the request headers**. Keys (names) are all lowercase.

💡 Use lowercase name to index header value.

```
local req_headers = ngx.req.get_headers()
setmetatable(req_headers, nil)
local user_agent = req_headers["user-agent"]
local cookie = req_headers["cookie"]
local x_forwarded_for = req_headers["x-forwarded-for"]
```

access log

```
http {  
    access_log logs/error.log main buffer=4096;  
}
```

- influenced by disk
- too many `write()` system calls (if buffer too small)

💡 flush logs to external services with `lua-resty-logger-socket`!

Good Programming Habits

- 💡 avoid overusing global variables
- 💡 avoid inefficient string concatenations
- 💡 avoid too much table resize
- 💡 use `lua-resty-core`
- 💡 use `JIT-compiled` functions
- 💡 use ffi to call your C functions

<http://wiki.luajit.org/NYI>

<https://blog.codingnow.com/cloud/LuaTips>

Name Resolution

- ◉ Misunderstanding about Nginx name resolution
- ◉ Nginx runtime DNS resolver's drawbacks
- ◉ Good solution based on `ngx_lua`

```
upstream backend {  
    server upyun.com weight=10 max_fails=30 fail_timeout=30s;  
}  
  
server {  
    listen *:8080;  
    location / {  
        proxy_pass http://backend;  
    }  
}
```

`upyun.com` will be resolved **only once** (while parsing the **server** directive).

```
server {  
    listen *:8080;  
    location / {  
        # without an explicit upstream group config.  
        proxy_pass http://upyun.com;  
    }  
}
```

Still `upyun.com` will be resolved `only once`.

```
server {  
    listen *:8080;  
    resolver 1.1.1.1:53 2.2.2.2:53 3.3.3.3:53 ipv6=off;  
    location / {  
        set $backend upyun.com;  
        proxy_pass http://$backend;  
    }  
}
```

Force ngx_proxy module to parse the “complex value”.

Now **upyun.com** will be resolved every time the ngx_proxy module runs.

Nginx runtime DNS resolver

- √ load balancing (RR)
- √ cache
- × backup
- × failover
- × use stale data

lua-resty-domain

- Based on lua-resty-dns.
- Combined with lua-resty-shcache, lua-resty-checkups.

- ✓ load balancing
- ✓ heartbeat and failover
- ✓ backend
- ✓ cache
- ✓ use stale data



```
_M.dns = {  
    enable = true,  
  
    cluster = {  
        {  
            servers = {  
                { host = "1.1.1.1", port = 53 },  
                { host = "2.2.2.2", port = 53 },  
            },  
        },  
        { -- backup  
            servers = {  
                { host = "3.3.3.3", port = 53 },  
            },  
        },  
    },  
  
    retrans = 2,  
    timeout = 1000,  
  
    -- shcache ttl parameters  
    positive_ttl = 15,  
    negative_ttl = 5,  
}
```

Patches for Cosocket

```
_G.ngx.socket.tcp = function(...)  
    local sock = tcp_sock(...)  
    sock.connect = function(sock, host, port, opts)  
        local addr = domain.toip(host)  
        return raw_connect(sock, addr, port, opts)  
    end  
    return sock  
end
```

Force Cosocket use our lua-resty-domain rather than the Nginx runtime DNS resolver, borrowed from Kong. :)

SSL Acceleration

Offload the RSA/ECDHE/SHA/... tasks to hardware (Intel QAT card), to reduce application's load.

Acceleration is not really “speed up”, but aims at increasing the application's throughput.

OpenResty + Asynchronous OpenSSL + Intel QAT

```
#openssl speed -elapsed rsa2048
```

You have chosen to measure elapsed time instead of user CPU time.

Doing 2048 bits private rsa's for 10s: 16640 2048 bits private RSA's in 10.00s

Doing 2048 bits public rsa's for 10s: 332804 2048 bits public RSA's in 10.00s

	sign	verify	sign/s	verify/s
rsa 2048 bits	0.000601s	0.000030s	1664.0	33280.4

```
#openssl speed -elapsed -engine qat -async_jobs 72 rsa2048
```


engine "qat" set.

You have chosen to measure elapsed time instead of user CPU time.

Doing 2048 bits private rsa's for 10s: 179738 2048 bits private RSA's in 10.01s

Doing 2048 bits public rsa's for 10s: 2213458 2048 bits public RSA's in 10.00s

	sign	verify	sign/s	verify/s
rsa 2048 bits	0.000056s	0.000005s	17955.8	221345.8



```
without-ssl-acceleration# ss -tn sport = :443 | wc -l  
32863  
with-ssl-acceleration# ss -tn sport = :443 | wc -l  
34689
```

CPU util ↓ ~10%。

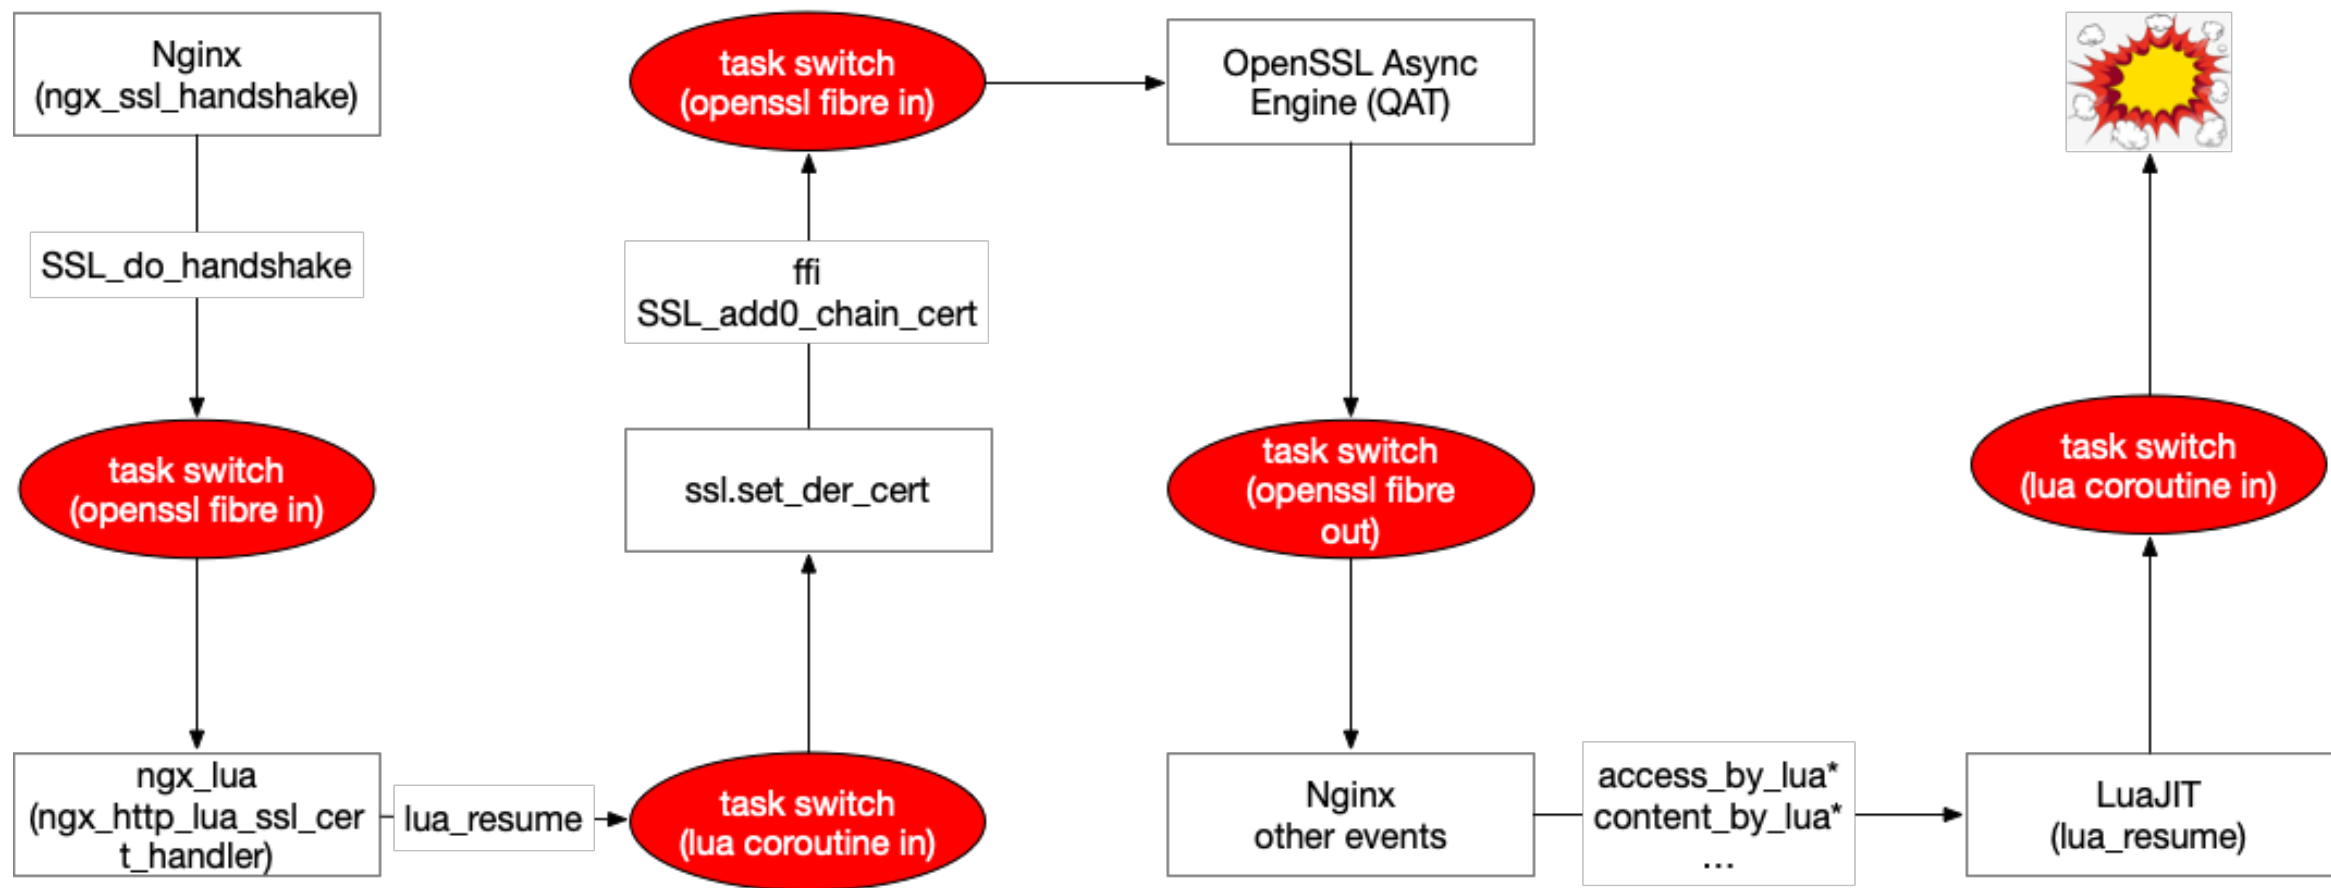
Conflict with LuaJIT

```
(gdb) bt
#0  0x0000000000000000 in ?? ()
#1  0x000000000047f8f1 in ngx_http_request_handler (ev=0x421ee98) at src/http/nginx_http_request.c:2259
#2  0x00000000004515d0 in ngx_event_process_posted (cycle=0x2c1c050, posted=0x9baa40 <ngx_posted_events>) at
src/event/nginx_event_posted.c:33
#3  0x000000000044f2ee in ngx_process_events_and_timers (cycle=0x2c1c050) at src/event/nginx_event.c:265
#4  0x000000000045c42b in ngx_worker_process_cycle (cycle=0x2c1c050, data=0x5) at src/os/unix/
ngx_process_cycle.c:792
#5  0x0000000000458eda in ngx_spawn_process (cycle=0x2c1c050, proc=0x45c37d <ngx_worker_process_cycle>,
data=0x5, name=0x6afd70 "worker process", respawn=-4)
    at src/os/unix/nginx_process.c:202
#6  0x000000000045b570 in ngx_start_worker_processes (cycle=0x2c1c050, n=16, type=-4) at src/os/unix/
ngx_process_cycle.c:373
#7  0x000000000045b0cd in ngx_master_process_cycle (cycle=0x2c1c050) at src/os/unix/nginx_process_cycle.c:256
#8  0x000000000041cef0 in main (argc=1, argv=0x7fff66d95838) at src/core/nginx.c:393
```

```
(gdb) bt
#0  0x00007f0c1a6aade2 in lj_vm_growstack_f () from /usr/local/marco/luajit/lib/libluajit-5.1.so.2
#1  0x0000000000550c13 in ngx_http_lua_run_thread (L=0x41a00378, r=0x1767f80, ctx=0x13793f0, nrets=0)
    at /disk/ssd2/alex_workflow/marco/deps/lua-nginx-module-0.10.11h/src/nginx_http_lua_util.c:1013
#2  0x000000000057825c in ngx_http_lua_ssl_cert_by_chunk (L=0x41a00378, r=0x1767f80)
    at /disk/ssd2/alex_workflow/marco/deps/lua-nginx-module-0.10.11h/src/nginx_http_lua_ssl_certby.c:527
#3  0x0000000000577457 in ngx_http_lua_ssl_cert_handler_file (r=0x1767f80, lscf=0x12c2138, L=0x41a00378)
    at /disk/ssd2/alex_workflow/marco/deps/lua-nginx-module-0.10.11h/src/nginx_http_lua_ssl_certby.c:57
#4  0x0000000000577c2c in ngx_http_lua_ssl_cert_handler (ssl_conn=0x1765790, data=0x0)
    at /disk/ssd2/alex_workflow/marco/deps/lua-nginx-module-0.10.11h/src/nginx_http_lua_ssl_certby.c:315
#5  0x00007f0c1a1e544a in tls_post_process_client_hello (s=0x1765790, wst=WORK_MORE_B) at ssl/statem/sta
#6  0x00007f0c1a1e2d2f in ossl_statem_server_post_process_message (s=0x1765790, wst=WORK_MORE_A) at ssl/
#7  0x00007f0c1a1cfe52 in read_state_machine (s=0x1765790) at ssl/statem/statem.c:660
#8  0x00007f0c1a1cf7a9 in state_machine (s=0x1765790, server=1) at ssl/statem/statem.c:428
#9  0x00007f0c1a1cf33b in ossl_statem_accept (s=0x1765790) at ssl/statem/statem.c:251
#10 0x00007f0c1a1b642d in ssl_do_handshake_intern (vargs=0x12c6730) at ssl/ssl_lib.c:3467
#11 0x00007f0c19d0770f in async_start_func () at crypto/async/async.c:154
#12 0x00007f0c199108f0 in __malloc_info (fp=0x7ffea8f29be0, options=<optimized out>) at malloc.c:5196
#13 0x00000000001f6b870 in ?? ()
#14 0x0000000000000000 in ?? ()
```

The devil is

- ◉ JIT compiler?
- ◉ Asynchronous OpenSSL mode?
- ◉ Lua task switch (Light Thread) in `ssl_certificate_by_lua*`?
- ◉ Nginx self?
- ◉ blah blah blah ...



System Stack -> OpenSSL Stack -> Lua Stack -> OpenSSL Stack -> System Stack -> Lua Stack

Reentry LuaJIT without calling `lua_yield()`!

Fixup

Prohibit OpenSSL fibres' task switch in **SSL-relevant APIs**.

- ssl.set_der_cert: ngx_http_lua_ffi_ssl_set_der_certificate
- ssl.set_der_priv_key: ngx_http_lua_ffi_ssl_set_der_private_key
- ssl.validate_ocsp_response: ngx_http_lua_ffi_ssl_validate_ocsp_response

```
#if OPENSSL_VERSION_NUMBER >= 0x10100000L
    ASYNC_block_pause();
#endif
.....

#if OPENSSL_VERSION_NUMBER >= 0x10100000L
    ASYNC_unblock_pause();
#endif
```

Thanks !

